

Product Requirements Document: Legal Intake & Triage Platform

Executive Summary

A conversational AI-powered legal intake system that adapts to user confidence levels, providing both fast-track request submission for experienced users and guided discovery for those uncertain about their legal needs. The platform reduces friction in accessing legal support while improving triage accuracy and enabling self-service resolution where appropriate.

Problem Statement

Current State:

- Internal teams struggle to determine if/when they need legal review
- Traditional intake forms create high friction and incomplete submissions
- Legal teams receive misrouted requests and lack context for proper triage
- Users fear "bothering" legal with questions, leading to either avoidance or last-minute urgent requests

Impact:

- Delayed legal reviews affecting business timelines
 - Legal team overwhelmed with low-complexity or unnecessary requests
 - Risk exposure from teams proceeding without proper legal guidance
 - Poor user experience leading to workarounds and shadow processes
-

Goals & Success Metrics

Primary Goals

1. Reduce time-to-appropriate-legal-resource by 50%
2. Increase self-service resolution rate by 30%
3. Improve initial request completeness (reduce back-and-forth by 40%)
4. Enable accurate auto-routing for 80% of requests

Success Metrics

Efficiency Metrics:

- Average time to submit request: < 3 minutes
- Requests requiring clarification: < 20%
- Auto-routing accuracy: > 80%
- Self-service resolution rate: > 25% of initiated requests

User Experience Metrics:

- User satisfaction score (CSAT): > 4.5/5
- Repeat user engagement rate: > 60%
- "Unsure" flow completion rate: > 70%

Business Impact Metrics:

- Legal team time saved on triage: 10+ hours/week
 - Reduction in urgent/last-minute requests: 30%
 - Knowledge base article views: 3x increase
-

User Personas

Persona 1: Confident Chris (Marketing Manager)

- **Experience:** Has worked with legal team 10+ times
- **Behavior:** Knows exactly what they need; wants speed
- **Pain Point:** Lengthy forms waste time when they already know the drill
- **Need:** Fast-track submission with minimal clicks

Persona 2: Uncertain Uma (Product Designer)

- **Experience:** Rarely interacts with legal; unclear on boundaries
- **Behavior:** Hesitant, asks questions, fears looking uninformed
- **Pain Point:** Doesn't know if situation requires legal review
- **Need:** Guided discovery without judgment; educational content

Persona 3: First-Timer Fiona (New Employee)

- **Experience:** Never used legal intake before
 - **Behavior:** Needs onboarding and context
 - **Pain Point:** Doesn't understand process, timelines, or expectations
 - **Need:** Clear guidance and reassurance
-

Core Features & Requirements

1. Intelligent Entry Point

Feature: Single conversational interface that adapts to user confidence

Requirements:

- **FR-1.1:** Open-ended initial prompt: "What are you working on?" or "How can we help?"
- **FR-1.2:** NLP detection of confidence signals (specific terminology vs. uncertain language)
- **FR-1.3:** Dynamic routing to Direct Path or Guided Discovery based on input analysis
- **FR-1.4:** Optional explicit mode selection: "I know what I need" vs. "Help me figure this out"
- **FR-1.5:** Context awareness: Display user's team, past requests, common patterns

Acceptance Criteria:

- System correctly identifies confident vs. uncertain users in 85%+ of cases
 - Users can manually switch between modes at any point
 - Interface loads in < 2 seconds with personalized context
-

2. Direct Path (Fast-Track Mode)

Feature: Streamlined submission for users who know their request type

Requirements:

- **FR-2.1:** Visual category selection with icons:
 - Contract Review
 - Marketing/Advertising Review
 - Partnership/Vendor Questions
 - Employment/HR Matters
 - Regulatory/Compliance
 - Intellectual Property
 - Other Legal Question
- **FR-2.2:** Progressive disclosure: Only show fields relevant to selected category
- **FR-2.3:** Smart defaults based on user history and team context
- **FR-2.4:** File upload with drag-and-drop (support .docx, .pdf, .xlsx, Google Docs links)
- **FR-2.5:** Optional urgency indicator with context: "Why is this urgent?"
- **FR-2.6:** Pre-submission summary: "Here's what we'll send to [Attorney Name]"
- **FR-2.7:** Expected timeline display: "Typical turnaround: 2-3 business days"

Acceptance Criteria:

- 90% of direct path requests completed in < 3 minutes
 - < 15% of requests require follow-up for missing information
 - Users can save drafts and return later
-

3. Guided Discovery Mode

Feature: Conversational triage for uncertain users

Requirements:


Phase 1: Situation Understanding




- **FR-3.1:** Open-ended questions in natural language
 - "What's the business goal or project?"
 - "What's making you think this might need legal review?"
 - "Have you or your team done something similar before?"

Phase 2: Risk Detection

- **FR-3.2:** Contextual follow-up questions based on initial responses:
 - Data collection triggers → Privacy/GDPR questions
 - Third-party involvement → Contract/vendor questions
 - Marketing content → Advertising claims questions
 - Geographic scope → Regulatory jurisdiction questions
- **FR-3.3:** Question branching logic based on detected risk factors
- **FR-3.4:** Educational tooltips explaining why each question matters

Phase 3: Intelligent Outcome

- **FR-3.5:** AI-generated assessment with one of four outcomes:
 -  **Needs Legal Review:** Auto-create request with collected context

-  **Might Need Review:** Provide decision checklist + option to submit
-  **Likely Fine:** Best practices + relevant policies to follow
-  **Self-Service:** Link to knowledge base articles, templates, FAQs
- **FR-3.6:** Transparent reasoning: "Based on [X, Y, Z], here's why..."
- **FR-3.7:** Always offer escalation: "Still want to talk to someone? Create request"

Acceptance Criteria:

- 70%+ completion rate for guided discovery flow
- 30%+ self-service resolution rate
- User confidence rating (post-interaction survey) > 4/5

4. AI-Powered Intelligence

Feature: Smart assistance throughout the experience

Requirements:

- **FR-4.1:** Auto-detection of urgency from language cues ("launching tomorrow," "deadline," "urgent")
- **FR-4.2:** Contextual pre-fill from connected systems:
 - Pull draft documents from Google Drive/SharePoint
 - Reference related Slack conversations
 - Link to similar past requests
- **FR-4.3:** Proactive self-service suggestions:
 - "Based on similar questions, here's our standard NDA template"
 - "Check if this FAQ answers your question: [link]"
- **FR-4.4:** Smart routing logic:
 - Match request to attorney expertise areas
 - Consider current workload and availability
 - Flag conflicts of interest or specialized needs
- **FR-4.5:** Auto-generated request summaries for legal team review
- **FR-4.6:** Learning system: Improve routing/triage based on attorney feedback

Acceptance Criteria:

- 80%+ auto-routing accuracy (validated by legal team)
- 25%+ of users engage with suggested self-service resources
- < 10% of auto-generated summaries require significant manual editing

5. Knowledge Base Integration

Feature: Contextual self-service resources

Requirements:

- **FR-5.1:** Searchable knowledge base with:
 - Legal policies and guidelines
 - Standard contract templates
 - Process documentation
 - FAQ by topic area
- **FR-5.2:** Inline article suggestions during triage
- **FR-5.3:** "Was this helpful?" feedback on all articles

- **FR-5.4:** Most-viewed and recently-updated content surfaced
- **FR-5.5:** Role-based content recommendations (Marketing sees advertising guidelines)

Acceptance Criteria:

- Knowledge base covers 80% of common request types
 - 3x increase in knowledge base engagement vs. current state
 - 40%+ of article viewers do not proceed to create request
-

6. Transparency & Status Tracking

Feature: Clear expectations and real-time updates

Requirements:

- **FR-6.1:** Immediate confirmation with:
 - Assigned attorney name and photo
 - Expected response timeline
 - Request reference number
- **FR-6.2:** Status tracking dashboard showing:
 - All pending requests
 - Current status (Submitted → Triaged → In Review → Complete)
 - Estimated completion date
- **FR-6.3:** Proactive notifications:
 - Email/Slack when request is assigned
 - Updates when status changes
 - Reminder if additional info needed
- **FR-6.4:** SLA visibility:
 - "This request type typically takes 2-3 days"
 - "Your request is on track / slightly delayed"

Acceptance Criteria:

- 90%+ of requests meet stated SLA expectations
 - < 5% of users contact legal directly asking for status updates
 - User satisfaction with transparency: > 4.5/5
-

7. Legal Team Interface

Feature: Efficient triage and case management for attorneys

Requirements:

- **FR-7.1:** Unified inbox with:
 - Filterable by urgency, type, attorney assignment
 - Rich request cards showing key context at a glance
 - Bulk actions (reassign, update status, close)
- **FR-7.2:** Request detail view with:
 - AI-generated summary at top
 - Full conversation history
 - Attached files and links
 - Related past requests from same user/team
- **FR-7.3:** Quick actions:
 - Accept/reassign request

- Request more information (templated responses)
- Mark complete with resolution notes
- Escalate to senior attorney
- **FR-7.4:** Feedback mechanisms:
 - Correct AI routing/triage decisions
 - Flag incomplete requests for system learning
 - Rate request quality
- **FR-7.5:** Analytics dashboard:
 - Request volume trends by type/team
 - Average resolution time by category
 - Attorney workload balancing
 - Self-service deflection rate

Acceptance Criteria:

- Legal team spends < 5 minutes on average per triage decision
 - 90%+ of requests have sufficient context for initial assessment
 - Attorney satisfaction with platform: > 4/5
-

User Flows

Flow 1: Confident User - Contract Review

1. User lands on platform → sees "What are you working on?"
2. Types: "Need standard NDA reviewed for new vendor"
3. AI detects confidence → shows fast-track categories
4. User clicks "Contract Review"
5. Selects "NDA" from contract type dropdown
6. Uploads document + adds vendor name "Acme Corp"
7. Optional: Sets target signature date
8. Reviews summary: "Sending to Sarah (Contracts) - typically 2 days"
9. Submits → receives confirmation with tracking link

Time to complete: 2 minutes

Flow 2: Uncertain User - Marketing Campaign

1. User lands on platform → sees "What are you working on?"
2. Types: "Planning an Instagram contest but not sure about rules"
3. AI detects uncertainty → enters guided mode
4. AI asks: "Tell me about the contest - what can people win?"
5. User: "Free product, \$500 value, need to tag our account"
6. AI asks: "Will you collect any personal information beyond Instagram handles?"
7. User: "We'll ask for email to ship the product"
8. AI asks: "What countries can people enter from?"
9. User: "US and Canada"
10. AI assessment: ⚠️ "This involves sweepstakes regulations, data collection, and multi-jurisdiction rules - you should get legal review"
11. Shows: Quick guide to what legal will need (official rules, privacy notice, etc.)
12. User clicks: "Yes, create request for me"
13. AI pre-fills request with conversation context → submits

Time to complete: 4 minutes

Flow 3: Self-Service Resolution

1. User lands on platform → sees "What are you working on?"
2. Types: "Can I use competitor name in our blog post?"
3. AI suggests: "This might be covered in our content guidelines - want to check first?"
4. User clicks: "Yes, show me"
5. Opens article: "Referential Use of Trademarks in Marketing Content"
6. Article clearly answers question with examples
7. User rates: "This answered my question" 👍
8. Optional: "Still need help?" → Create request
9. User exits satisfied

Time to complete: 1 minute

Technical Requirements

Platform Architecture

TR-1: Frontend

- React-based web application
- Mobile-responsive design (support iOS/Android browsers)
- Accessibility compliant (WCAG 2.1 AA)
- Real-time updates via WebSocket

TR-2: Backend

- RESTful API architecture
- Authentication via SSO (SAML 2.0 / OAuth)
- Role-based access control (RBAC)

TR-3: AI/ML Components

- Natural language processing for intent detection
- Classification model for request routing (85%+ accuracy)
- Sentiment analysis for urgency detection
- Continuous learning pipeline with attorney feedback loop

TR-4: Integrations

- Google Workspace (Drive, Docs, Calendar)
- Microsoft 365 (SharePoint, OneDrive, Outlook)
- Slack / Microsoft Teams for notifications
- Existing legal case management system (if applicable)
- Analytics platform (e.g., Mixpanel, Amplitude)

TR-5: Data & Security

- End-to-end encryption for sensitive documents
- Data retention policy enforcement
- Audit logging for all actions

- GDPR/CCPA compliance
- SOC 2 Type II certification

TR-6: Performance

- Page load time: < 2 seconds (p95)
 - API response time: < 500ms (p95)
 - 99.9% uptime SLA
 - Support 500+ concurrent users
-

Data Requirements

Data Capture

User Input Data:

- Request description and context
- Category selection
- Uploaded documents and links
- Urgency indicators
- User profile (team, role, location)

System-Generated Data:

- Confidence score (certain vs. uncertain user)
- Risk factors detected
- Suggested resources
- Routing recommendations
- Timestamp metadata

Legal Team Data:

- Attorney assignments
- Status updates
- Resolution notes
- Time tracking
- Quality ratings

Analytics & Reporting

User Metrics:

- Requests by team/department
- Request types distribution
- Completion funnel (started vs. submitted)
- Self-service resolution rate
- User satisfaction scores

Operational Metrics:

- Average time to assignment
- Average time to resolution by category
- SLA compliance rate
- Attorney utilization and workload

- Most common request types

AI Performance Metrics:

- Routing accuracy (vs. attorney corrections)
 - Urgency detection accuracy
 - Knowledge base suggestion relevance
 - Model confidence scores
-

Design Requirements

Visual Design Principles

DR-1: Approachable, Not Intimidating

- Conversational tone throughout
- Friendly micro-copy ("No worries!" vs. "Error")
- Minimal legal jargon; tooltips when necessary
- Progress indicators to reduce anxiety

DR-2: Clear Visual Hierarchy

- Primary action always obvious
- Secondary options available but not distracting
- Consistent component library
- Generous white space

DR-3: Trust & Transparency

- Show who will handle request (photo + name)
- Explain why questions are asked
- Display expected timelines upfront
- Make AI reasoning visible

DR-4: Accessibility

- Keyboard navigation support
- Screen reader compatible
- Sufficient color contrast (4.5:1 minimum)
- Clear focus indicators
- Support for dyslexia-friendly fonts

Key UI Components

Conversational Interface:

- Chat-like message bubbles
- Typing indicators for AI responses
- Inline question cards with radio/checkbox options
- Rich media support (file previews, embedded videos)

Category Selection:

- Large clickable cards with icons

- 2-3 word labels + brief description on hover
- Visual distinction between categories
- Search/filter if > 8 categories

Status Dashboard:

- Card-based layout for each request
 - Color-coded status badges
 - Quick view of key details
 - Expandable for full history
-

Implementation Phases

Phase 1: MVP (Months 1-3)

Scope:

- Basic conversational intake with 5 core categories
- Direct path for confident users
- Simple guided discovery (5-7 questions max)
- Manual attorney assignment
- Email notifications
- Basic analytics dashboard

Success Criteria:

- 50+ requests processed through platform
 - 70%+ user satisfaction
 - 50% reduction in incomplete requests
-

Phase 2: AI Enhancement (Months 4-6)

Scope:

- AI-powered routing and assignment
- Urgency detection
- Knowledge base integration with inline suggestions
- Self-service outcome paths
- Slack/Teams integration
- Enhanced analytics

Success Criteria:

- 80%+ auto-routing accuracy
 - 20%+ self-service resolution rate
 - 10 hours/week legal team time saved
-

Phase 3: Advanced Intelligence (Months 7-9)

Scope:

- Context pre-fill from connected systems

- Predictive triage (proactive suggestions)
- Learning system with feedback loop
- Advanced workflow automation
- Custom reporting for leadership
- Mobile app (optional)

Success Criteria:

- 30%+ self-service resolution rate
 - < 2 minute average submission time
 - 90%+ repeat user adoption
-

Phase 4: Optimization & Scale (Months 10-12)

Scope:

- Multi-language support
- Advanced personalization
- Integration with contract lifecycle management
- AI-generated draft responses for legal team
- Predictive workload planning
- External vendor portal (optional)

Success Criteria:

- Platform handles 100% of legal intake
 - 85%+ overall user satisfaction
 - Measurable reduction in legal cycle times
-

Risk Assessment & Mitigation

Technical Risks

R-1: AI Routing Accuracy

- **Risk:** Low accuracy leads to misrouted requests, wasted time
- **Mitigation:** Start with human-in-the-loop validation; continuous learning; fallback to manual routing if confidence < 70%

R-2: Integration Complexity

- **Risk:** Connecting to multiple enterprise systems delays launch
- **Mitigation:** Phase integrations; start with email/basic auth; use webhooks where possible

R-3: Performance Under Load

- **Risk:** System slows during peak usage
- **Mitigation:** Load testing before launch; auto-scaling infrastructure; caching strategy

Adoption Risks

R-4: User Resistance to Change

- **Risk:** Teams continue using email/Slack for legal requests
- **Mitigation:** Executive sponsorship; team-by-team rollout with champions; showcase time savings; make platform easier than alternatives

R-5: Legal Team Adoption

- **Risk:** Attorneys prefer old intake methods
- **Mitigation:** Early involvement in design; highlight triage time savings; provide training; maintain hybrid option during transition

Content Risks

R-6: Knowledge Base Incompleteness

- **Risk:** AI suggests irrelevant or outdated resources
- **Mitigation:** Content audit before launch; ongoing content maintenance; track "not helpful" feedback; legal team review cycle

R-7: AI Liability/Advice Concerns

- **Risk:** Users interpret AI guidance as legal advice
 - **Mitigation:** Clear disclaimers; AI provides triage only, not legal opinions; escalation paths always available; attorney approval of all self-service outcomes
-

Dependencies

External Dependencies

- SSO provider configuration
- Access to Google Workspace / M365 APIs
- Slack/Teams app approval
- AI/ML platform licensing (e.g., OpenAI, Anthropic)
- Legal team availability for training data and feedback

Internal Dependencies

- Executive sponsorship and budget approval
- Legal team champions identified
- IT security review and approval
- Content team for knowledge base creation
- Change management resources for rollout

Team Requirements

- Product Manager (1 FTE)
 - Product Designer (0.5 FTE)
 - Frontend Engineers (2 FTE)
 - Backend Engineers (2 FTE)
 - ML/AI Engineer (1 FTE)
 - Legal Ops Specialist (0.5 FTE)
 - Content Writer (0.25 FTE)
-

Open Questions

1. **Scope:** Should we include external vendor/partner access in MVP or Phase 4?
 2. **Routing:** How do we handle requests requiring multiple legal specialties?
 3. **Urgency:** Should we auto-escalate certain keywords (e.g., "lawsuit," "subpoena") regardless of user input?
 4. **Data retention:** What's the policy for storing conversation history and uploaded documents?
 5. **Global teams:** Do we need multi-language support in MVP or can we start English-only?
 6. **Integration priority:** Which connected systems provide the most value - prioritize which integrations?
 7. **Self-service limits:** What types of questions should NEVER have self-service outcomes?
 8. **Metrics ownership:** Who owns ongoing platform analytics and optimization?
-

Appendix

Glossary

- **Triage:** Initial assessment and categorization of legal requests
- **Self-service resolution:** User finds answer without creating formal request
- **Direct path:** Fast-track flow for users who know their request type
- **Guided discovery:** Conversational flow for uncertain users
- **SLA (Service Level Agreement):** Expected timeframe for legal team response

References

- Current legal intake form (baseline for comparison)
 - Legal team interview notes (pain points and requirements)
 - User research findings (20 interviews across departments)
 - Competitive analysis (LawVu, Ironclad, Onit)
 - Industry benchmarks (legal ops metrics)
-

Document Version: 1.0
Last Updated: October 21, 2025
Owner: [Product Manager Name]
Reviewers: Legal Operations Lead, Engineering Lead, Design Lead
Status: Draft for Review